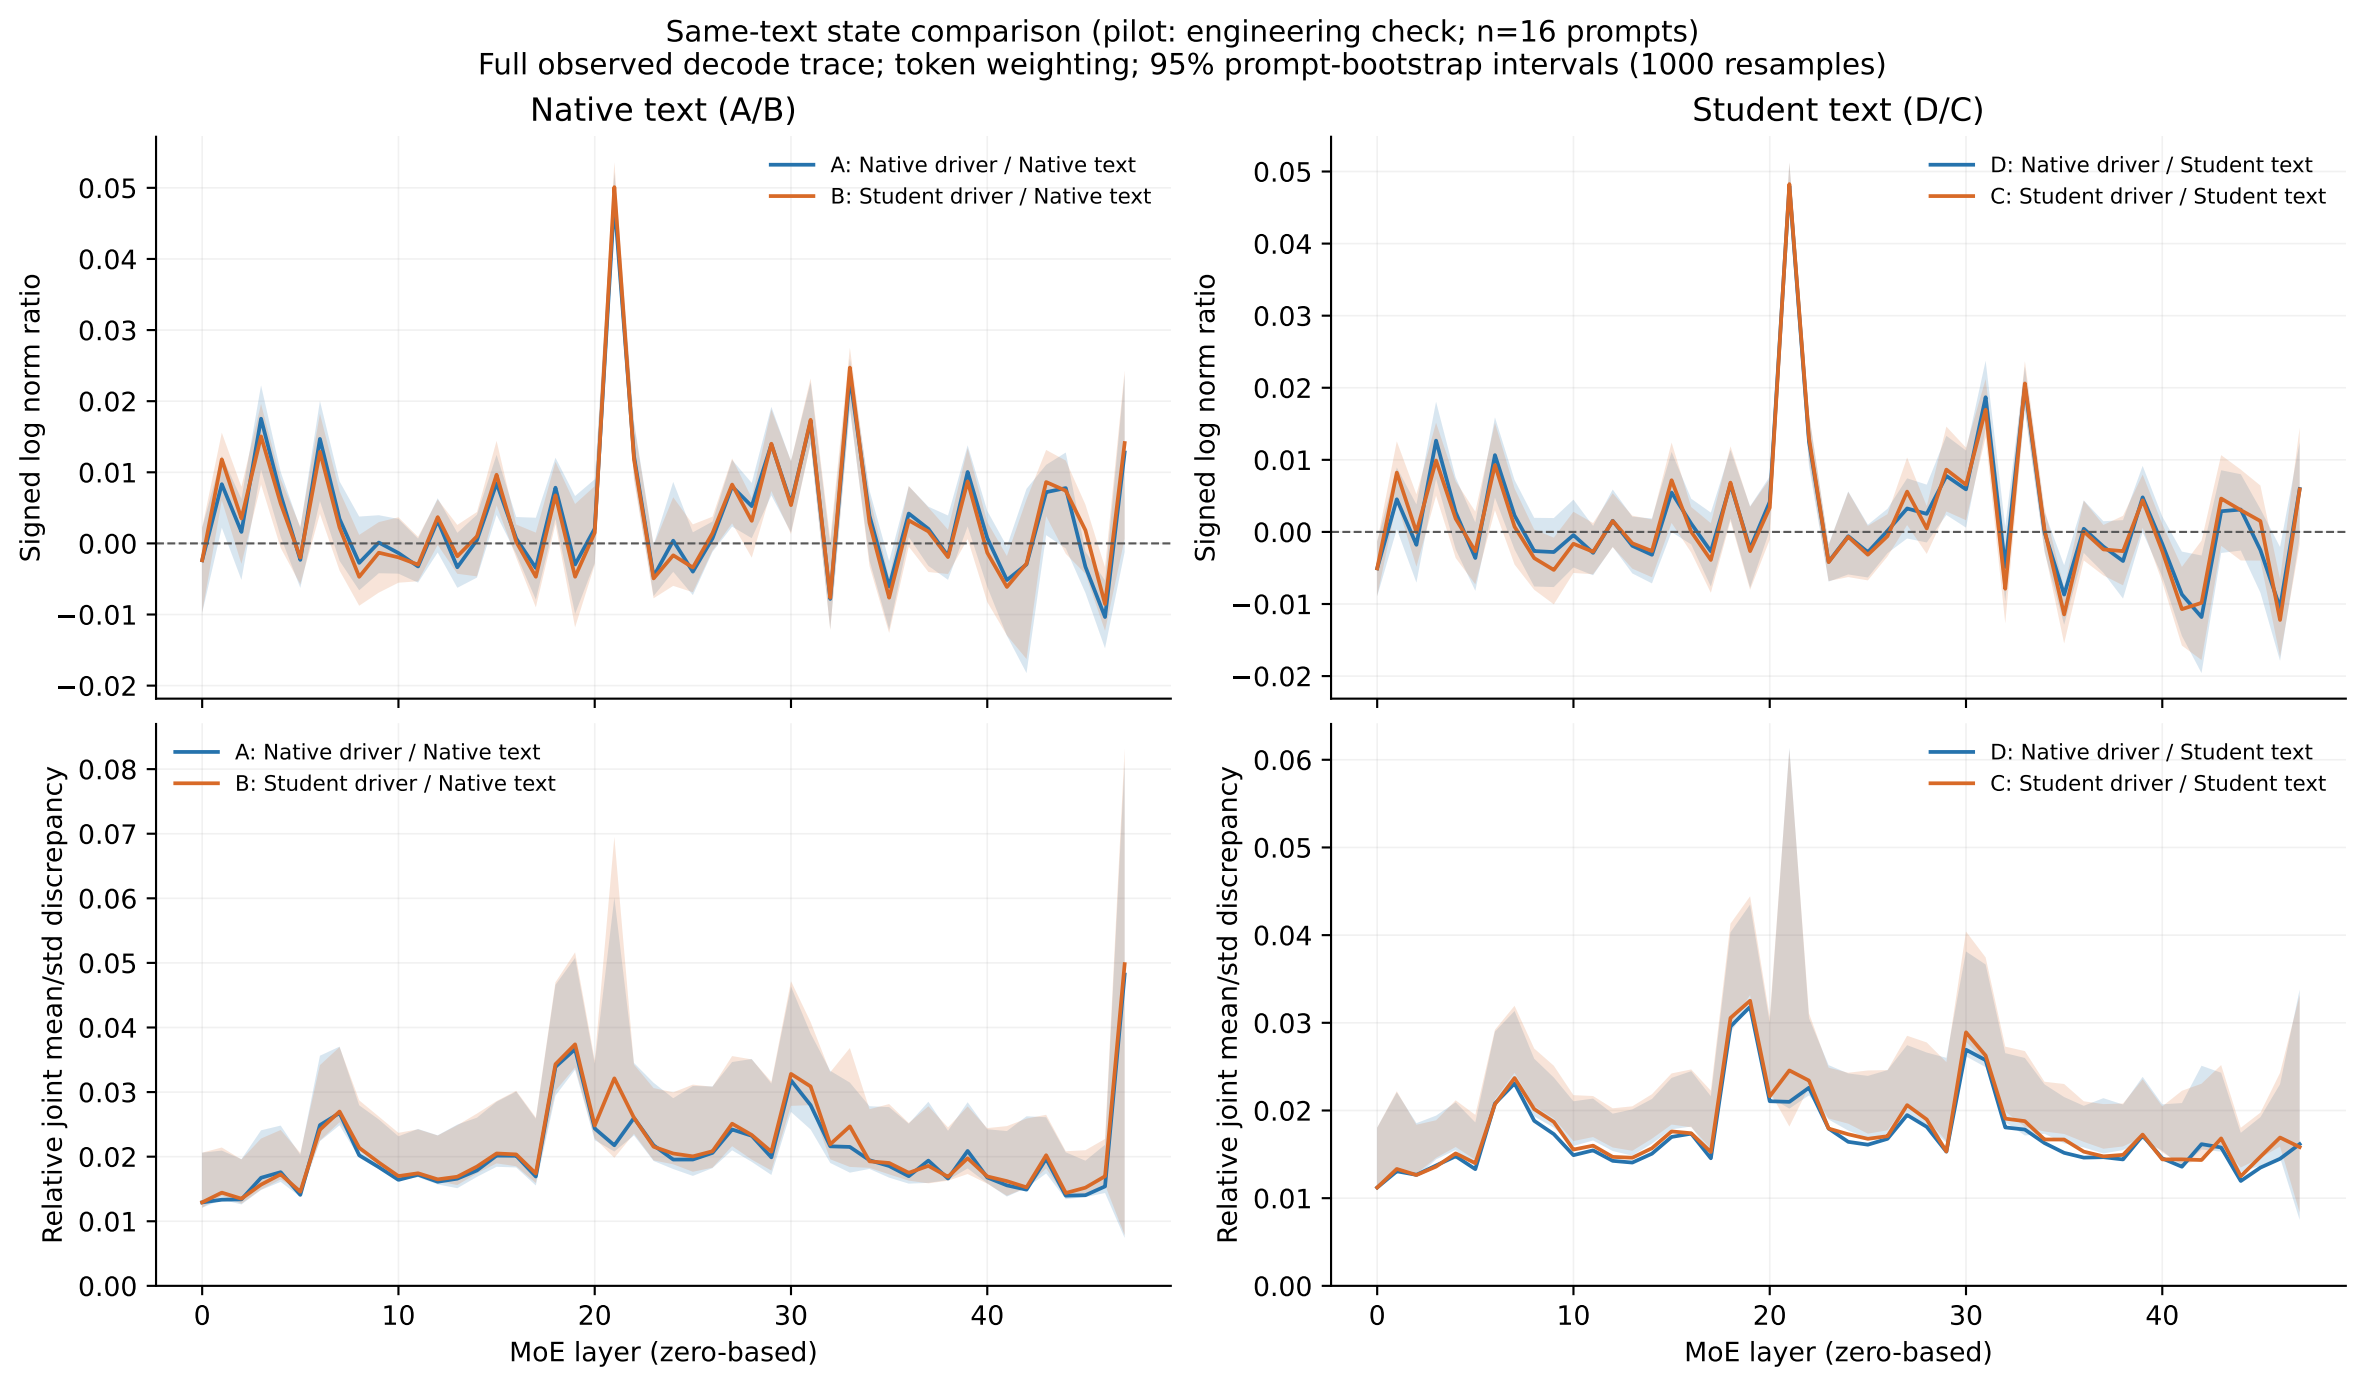
Same-text state comparison (pilot: engineering check; n=16 prompts)
Full observed decode trace; token weighting; 95% prompt-bootstrap intervals (1000 resamples)
Native text (A/B)
Student text (D/C)
A: Native driver / Native text
B: Student driver / Native text
D: Native driver / Student text
C: Student driver / Student text
Signed log norm ratio
Relative joint mean/std discrepancy
MoE layer (zero-based)